

# 1. Import Libraries

In [36]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from xgboost import XGBClassifier
RANDOM_STATE = 55 # for reproducibility
```

# 2. Data Exploration

In [2]:

```
# Read the excel file and check first few lines
pumpkin = pd.read_excel("/Users/xinyizhang/Desktop/Weill Cornell Medicine/自学
pumpkin.head()
```

Out[2]:

	Area	Perimeter	Major_Axis_Length	Minor_Axis_Length	Convex_Area	Equiv_Diameter	Eccentricity
0	56276	888.242	326.1485	220.2388	56831	267.6805	0.3123614
1	76631	1068.146	417.1932	234.2289	77280	312.3614	0.3019822
2	71623	1082.987	435.8328	211.0457	72663	290.8899	0.2901207
3	66458	992.051	381.5638	222.5322	67118	290.8899	0.2890555
4	66107	998.146	383.8883	220.4545	67117	290.1207	0.2880891

In [3]:

```
pumpkin.info()
pumpkin.shape
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 13 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Area             2500 non-null    int64  
 1   Perimeter        2500 non-null    float64 
 2   Major_Axis_Length 2500 non-null    float64 
 3   Minor_Axis_Length 2500 non-null    float64 
 4   Convex_Area      2500 non-null    int64  
 5   Equiv_Diameter   2500 non-null    float64 
 6   Eccentricity     2500 non-null    float64 
 7   Solidity          2500 non-null    float64 
 8   Extent            2500 non-null    float64 
 9   Roundness          2500 non-null    float64 
 10  Aspect_Ration     2500 non-null    float64 
 11  Compactness        2500 non-null    float64 
 12  Class              2500 non-null    object 
dtypes: float64(10), int64(2), object(1)
memory usage: 254.0+ KB
```

Out[3]: (2500, 13)

The data has 2500 rows and 13 columns, among which the first 12 columns are features and the last column is the target variable

In [4]:

```
# Summarize the unique labels in the Class column
class_summary = pumpkin['Class'].value_counts()
class_summary
```

Out[4]:

Class	Count
Çerçeveilik	1300
Ürgüp Sivrisi	1200
Name: count, dtype: int64	

In [5]:

```
# Summarize the numeric pumpkin data
pumpkin_summary = pumpkin.describe()
pumpkin_summary
```

Out[5]:

	Area	Perimeter	Major_Axis_Length	Minor_Axis_Length	Convex_Area	E
<b>count</b>	2500.000000	2500.000000	2500.000000	2500.000000	2500.000000	
<b>mean</b>	80658.220800	1130.279015	456.601840	225.794921	81508.084400	
<b>std</b>	13664.510228	109.256418	56.235704	23.297245	13764.092788	
<b>min</b>	47939.000000	868.485000	320.844600	152.171800	48366.000000	
<b>25%</b>	70765.000000	1048.829750	414.957850	211.245925	71512.000000	
<b>50%</b>	79076.000000	1123.672000	449.496600	224.703100	79872.000000	
<b>75%</b>	89757.500000	1203.340500	492.737650	240.672875	90797.750000	
<b>max</b>	136574.000000	1559.450000	661.911300	305.818000	138384.000000	

In [6]:

```
# Describe the data by unique pumpkin seeds classes
class_description = pumpkin.groupby('Class').describe()
class_description
```

Out[6]:

	count	mean	std	min	25%	50%	75%	ma
<b>Class</b>								
<b>Çerçeveilik</b>	1300.0	78423.154615	11246.499728	55811.0	69777.75	76718.5	86277.75	107476.
<b>Ürgüp Sivrisi</b>	1200.0	83079.542500	15519.323847	47939.0	72482.50	81657.0	93815.75	136574.

2 rows × 96 columns


## Visualization

In [7]:

```
sns.scatterplot(data = pumpkin, x = 'Area', y = 'Perimeter', hue = 'Class')
```


Out[7]:

```
<Axes: xlabel='Area', ylabel='Perimeter'>
```




```
In [8]: sns.scatterplot(data = pumpkin, x = 'Major_Axis_Length', y = 'Minor_Axis_Leng
```

```
Out[8]: <Axes: xlabel='Major_Axis_Length', ylabel='Minor_Axis_Length'>
```



```
In [9]: sns.scatterplot(data = pumpkin, x = 'Roundness', y = 'Compactness', hue = 'Cl
```

Out[9]:



### 3. Random Forest Building

```
In [10]: features = [x for x in pumpkin.columns if x not in 'Class'] # Removing our ta
```

Split the data into training and test set

```
In [11]: X_train, X_test, y_train, y_test = train_test_split(pumpkin[features], pumpki
```

```
In [12]: print(f'Train samples: {len(X_train)})')
print(f'Validation samples: {len(X_test)})')
```

Train samples: 2000  
Validation samples: 500

Choose the most suitable hyperparameter

```
In [13]: # Minimum samples per leaf
min_samples_split_list = [2, 5, 10, 20, 50, 100, 250, 500, 700]

# Maximum depth of the tree
max_depth_list = [2, 4, 8, 16, 32, 64, None]

# Number of trees in the forest
n_list = [10, 50, 100, 200, 500]
```

In [14]:


```

accuracy_list_train = []
accuracy_list_test = []
for min_samples_split in min_samples_split_list:
    model = RandomForestClassifier(min_samples_split = min_samples_split,
                                    random_state = RANDOM_STATE).fit(X_train,y)
    predictions_train = model.predict(X_train) # The predicted values for the training set
    predictions_test = model.predict(X_test) # The predicted values for the test set
    accuracy_train = accuracy_score(predictions_train,y_train)
    accuracy_test = accuracy_score(predictions_test,y_test)
    accuracy_list_train.append(accuracy_train)
    accuracy_list_test.append(accuracy_test)

plt.title('Train x Validation metrics')
plt.xlabel('Minimum Samples Split')
plt.ylabel('Accuracy')
plt.xticks(ticks = range(len(min_samples_split_list))),labels=min_samples_split_list)
plt.plot(accuracy_list_train)
plt.plot(accuracy_list_test)
plt.legend(['Train','Validation'])

```

Out[14]: &lt;matplotlib.legend.Legend at 0x17dc6cf40&gt;



In [15]:


```

accuracy_list_train = []
accuracy_list_test = []
for max_depth in max_depth_list:
    model = RandomForestClassifier(max_depth = max_depth,
                                    random_state = RANDOM_STATE).fit(X_train,y)
    predictions_train = model.predict(X_train) # The predicted values for the training set
    predictions_test = model.predict(X_test) # The predicted values for the test set
    accuracy_train = accuracy_score(predictions_train,y_train)
    accuracy_test = accuracy_score(predictions_test,y_test)
    accuracy_list_train.append(accuracy_train)
    accuracy_list_test.append(accuracy_test)

```

```
plt.title('Train x Validation metrics')
plt.xlabel('Maximum Depth')
plt.ylabel('Accuracy')
plt.xticks(ticks = range(len(max_depth_list )),labels=max_depth_list)
plt.plot(accuracy_list_train)
plt.plot(accuracy_list_test)
plt.legend(['Train','Validation'])
```

Out[15]: &lt;matplotlib.legend.Legend at 0x17421c310&gt;




In [16]:

```
accuracy_list_train = []
accuracy_list_test = []
for n_estimators in n_list:
    model = RandomForestClassifier(n_estimators = n_estimators,
                                    random_state = RANDOM_STATE).fit(X_train,y
    predictions_train = model.predict(X_train) # The predicted values for the training set
    predictions_test = model.predict(X_test) # The predicted values for the test set
    accuracy_train = accuracy_score(predictions_train,y_train)
    accuracy_test = accuracy_score(predictions_test,y_test)
    accuracy_list_train.append(accuracy_train)
    accuracy_list_test.append(accuracy_test)

plt.title('Train x Validation metrics')
plt.xlabel('Number of Trees')
plt.ylabel('Accuracy')
plt.xticks(ticks = range(len(n_list)),labels=n_list)
plt.plot(accuracy_list_train)
plt.plot(accuracy_list_test)
plt.legend(['Train','Validation'])
```

Out[16]: &lt;matplotlib.legend.Legend at 0x17fcb5a60&gt;



In order to achieve the best result, choose the following numbers as hyperparamters:  
 Maximum Sample Split: 20 Only try to split a node if there are at least 20 samples in it  
 Maximum Depth: 8 Number of Trees: 100

## Final Random Forest Model

```
In [17]: random_forest_model = RandomForestClassifier(n_estimators = 100,
                                                 max_depth = 8,
                                                 min_samples_split = 20).fit(X_tr
```

## Evaluation

### Classification Report


```
In [18]: y_test_pred = random_forest_model.predict(X_test)
print(classification_report(y_test, y_test_pred))
```

	precision	recall	f1-score	support
Çerçeveilik	0.88	0.92	0.90	251
Ürgüp Sivrisi	0.92	0.87	0.89	249
accuracy			0.90	500
macro avg	0.90	0.90	0.90	500
weighted avg	0.90	0.90	0.90	500

## Confusion Matrix

In [19]:

```
class_labels = pumpkin['Class'].unique()
cm = confusion_matrix(y_test, y_test_pred)
sns.heatmap(cm, annot = True, fmt = 'd', cmap = 'Blues', xticklabels = class_
plt.title('Pumpkin Seed Confusion Matrix')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```




## Importance

In [43]:

```
importances = random_forest_model.feature_importances_
feat_importances = pd.Series(importances, index = features)
feat_importances.sort_values(ascending = False).plot(kind = 'bar')
plt.title("Feature Importance")
plt.show()
```

## Feature Importance



## 4. XGBoost

The boosting methods train several trees, but instead of them being uncorrelated to each other, now the trees are fit one after the other in order to minimize the error.

The model has the same parameters as a decision tree, plus the learning rate.

- The learning rate is the size of the step on the Gradient Descent method that the XGBoost uses internally to minimize the error on each train step.

### Split the data into training and testing set

```
In [20]: X_train, X_test, y_train, y_test = train_test_split(pumpkin[features], pumpkin['Type'], test_size=0.2, random_state=42)
```

```
In [21]: # Encode labels
label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train)
y_test_encoded = label_encoder.transform(y_test)
```

### Choose the most suitable hyperparameters

In [45]:

```
# Define the learning rate
learning_rate = [0.1, 0.2, 0.3, 0.5, 0.7, 1.0]

# Define the maximum depth of the tree
max_depth_list = [2, 4, 8, 16, 32, 64, None]
```

In [47]:

```
accuracy_list_test_xgb = []
accuracy_list_train_xgb = []
for lr in learning_rate:
    # Number of trees can be set to a higher value for XGBoost since model will
    # converge faster with more trees
    gb_model = XGBClassifier(n_estimators = 500, learning_rate = lr, verbose=False)
    gb_model.fit(X_train, y_train_encoded, eval_set=[(X_test, y_test_encoded)])
    prediction_train = gb_model.predict(X_train)
    prediction_test = gb_model.predict(X_test)
    accuracy_train = accuracy_score(prediction_train, y_train_encoded)
    accuracy_test = accuracy_score(prediction_test, y_test_encoded)
    accuracy_list_train_xgb.append(accuracy_train)
    accuracy_list_test_xgb.append(accuracy_test)


plt.title('Learning Rate vs Accuracy')
plt.xlabel('Learning Rate')
plt.ylabel('Accuracy')
plt.xticks(ticks=range(len(learning_rate)), labels = learning_rate)
plt.plot(accuracy_list_train_xgb)
plt.plot(accuracy_list_test_xgb)
plt.legend(['Train', "Validation"])
plt.show()
```

```
[0] validation_0-logloss:0.52957
[1] validation_0-logloss:0.43944
[2] validation_0-logloss:0.38977
[3] validation_0-logloss:0.35359
[4] validation_0-logloss:0.32719
[5] validation_0-logloss:0.31539
[6] validation_0-logloss:0.30277
[7] validation_0-logloss:0.29750
[8] validation_0-logloss:0.29406
[9] validation_0-logloss:0.29429
[10] validation_0-logloss:0.29286
[11] validation_0-logloss:0.29258
[12] validation_0-logloss:0.29295
[13] validation_0-logloss:0.29271
[14] validation_0-logloss:0.29479
[15] validation_0-logloss:0.29615
[16] validation_0-logloss:0.29605
[17] validation_0-logloss:0.29656
[18] validation_0-logloss:0.29605
[19] validation_0-logloss:0.29830
[20] validation_0-logloss:0.30036
[21] validation_0-logloss:0.30155
[0] validation_0-logloss:0.52957
[1] validation_0-logloss:0.43944
[2] validation_0-logloss:0.38977
[3] validation_0-logloss:0.35359
[4] validation_0-logloss:0.32719
[5] validation_0-logloss:0.31539
[6] validation_0-logloss:0.30277
[7] validation_0-logloss:0.29750
[8] validation_0-logloss:0.29406
[9] validation_0-logloss:0.29429
[10] validation_0-logloss:0.29286
```

```
[11] validation_0-logloss:0.29258
[12] validation_0-logloss:0.29295
[13] validation_0-logloss:0.29271
[14] validation_0-logloss:0.29479
[15] validation_0-logloss:0.29615
[16] validation_0-logloss:0.29605
[17] validation_0-logloss:0.29656
[18] validation_0-logloss:0.29605
[19] validation_0-logloss:0.29830
[20] validation_0-logloss:0.30036
[0] validation_0-logloss:0.52957
[1] validation_0-logloss:0.43944
[2] validation_0-logloss:0.38977
[3] validation_0-logloss:0.35359
[4] validation_0-logloss:0.32719
[5] validation_0-logloss:0.31539
[6] validation_0-logloss:0.30277
[7] validation_0-logloss:0.29750
[8] validation_0-logloss:0.29406
[9] validation_0-logloss:0.29429
[10] validation_0-logloss:0.29286
[11] validation_0-logloss:0.29258
[12] validation_0-logloss:0.29295
[13] validation_0-logloss:0.29271
[14] validation_0-logloss:0.29479
[15] validation_0-logloss:0.29615
[16] validation_0-logloss:0.29605
[17] validation_0-logloss:0.29656
[18] validation_0-logloss:0.29605
[19] validation_0-logloss:0.29830
[20] validation_0-logloss:0.30036
[21] validation_0-logloss:0.30155
[0] validation_0-logloss:0.52957
[1] validation_0-logloss:0.43944
[2] validation_0-logloss:0.38977
[3] validation_0-logloss:0.35359
[4] validation_0-logloss:0.32719
[5] validation_0-logloss:0.31539
[6] validation_0-logloss:0.30277
[7] validation_0-logloss:0.29750
[8] validation_0-logloss:0.29406
[9] validation_0-logloss:0.29429
[10] validation_0-logloss:0.29286
[11] validation_0-logloss:0.29258
[12] validation_0-logloss:0.29295
[13] validation_0-logloss:0.29271
[14] validation_0-logloss:0.29479
[15] validation_0-logloss:0.29615
[16] validation_0-logloss:0.29605
[17] validation_0-logloss:0.29656
[18] validation_0-logloss:0.29605
[19] validation_0-logloss:0.29830
[20] validation_0-logloss:0.30036
[0] validation_0-logloss:0.52957
[1] validation_0-logloss:0.43944
[2] validation_0-logloss:0.38977
[3] validation_0-logloss:0.35359
[4] validation_0-logloss:0.32719
[5] validation_0-logloss:0.31539
[6] validation_0-logloss:0.30277
[7] validation_0-logloss:0.29750
[8] validation_0-logloss:0.29406
[9] validation_0-logloss:0.29429
[10] validation_0-logloss:0.29286
```

```
[11] validation_0-logloss:0.29258
[12] validation_0-logloss:0.29295
[13] validation_0-logloss:0.29271
[14] validation_0-logloss:0.29479
[15] validation_0-logloss:0.29615
[16] validation_0-logloss:0.29605
[17] validation_0-logloss:0.29656
[18] validation_0-logloss:0.29605
[19] validation_0-logloss:0.29830
[20] validation_0-logloss:0.30036
[0] validation_0-logloss:0.52957
[1] validation_0-logloss:0.43944
[2] validation_0-logloss:0.38977
[3] validation_0-logloss:0.35359
[4] validation_0-logloss:0.32719
[5] validation_0-logloss:0.31539
[6] validation_0-logloss:0.30277
[7] validation_0-logloss:0.29750
[8] validation_0-logloss:0.29406
[9] validation_0-logloss:0.29429
[10] validation_0-logloss:0.29286
[11] validation_0-logloss:0.29258
[12] validation_0-logloss:0.29295
[13] validation_0-logloss:0.29271
[14] validation_0-logloss:0.29479
[15] validation_0-logloss:0.29615
[16] validation_0-logloss:0.29605
[17] validation_0-logloss:0.29656
[18] validation_0-logloss:0.29605
[19] validation_0-logloss:0.29830
[20] validation_0-logloss:0.30036
```

Learning Rate vs Accuracy



In [46]:

```
accuracy_list_test_xgb = []
accuracy_list_train_xgb = []
for depth in max_depth_list:
```


```
# Number of trees can be set to a higher value for XGBoost since model wi
xgb_model = XGBClassifier(n_estimators = 500, max_depth = depth, verboxit
xgb_model.fit(X_train, y_train_encoded, eval_set=[(X_test, y_test_encoded)
prediction_train = xgb_model.predict(X_train)
prediction_test = xgb_model.predict(X_test)
accuracy_train = accuracy_score(prediction_train, y_train_encoded)
accuracy_test = accuracy_score(prediction_test, y_test_encoded)
accuracy_list_train_xgb.append(accuracy_train)
accuracy_list_test_xgb.append(accuracy_test)

plt.title('Learning Rate vs Accuracy')
plt.xlabel('Depth')
plt.ylabel('Accuracy')
plt.xticks(ticks=range(len(max_depth_list)), labels = max_depth_list)
plt.plot(accuracy_list_train_xgb)
plt.plot(accuracy_list_test_xgb)
plt.legend(['Train', "Validation"])
plt.show()
```

```
[0] validation_0-logloss:0.54215
[1] validation_0-logloss:0.45459
[2] validation_0-logloss:0.40414
[3] validation_0-logloss:0.37337
[4] validation_0-logloss:0.35276
[5] validation_0-logloss:0.33243
[6] validation_0-logloss:0.32235
[7] validation_0-logloss:0.31157
[8] validation_0-logloss:0.30627
[9] validation_0-logloss:0.30165
[10] validation_0-logloss:0.29612
[11] validation_0-logloss:0.29369
[12] validation_0-logloss:0.29106
[13] validation_0-logloss:0.28805
[14] validation_0-logloss:0.28599
[15] validation_0-logloss:0.28437
[16] validation_0-logloss:0.28431
[17] validation_0-logloss:0.28405
[18] validation_0-logloss:0.28320
[19] validation_0-logloss:0.28203
[20] validation_0-logloss:0.28238
[21] validation_0-logloss:0.28216
[22] validation_0-logloss:0.28221
[23] validation_0-logloss:0.28378
[24] validation_0-logloss:0.28327
[25] validation_0-logloss:0.28504
[26] validation_0-logloss:0.28550
[27] validation_0-logloss:0.28530
[28] validation_0-logloss:0.28395
[0] validation_0-logloss:0.53268
[1] validation_0-logloss:0.44039
[2] validation_0-logloss:0.38649
[3] validation_0-logloss:0.35107
[4] validation_0-logloss:0.33148
[5] validation_0-logloss:0.31687
[6] validation_0-logloss:0.31061
[7] validation_0-logloss:0.30357
[8] validation_0-logloss:0.29831
[9] validation_0-logloss:0.29613
[10] validation_0-logloss:0.29318
[11] validation_0-logloss:0.29221
[12] validation_0-logloss:0.28951
[13] validation_0-logloss:0.28946
[14] validation_0-logloss:0.29206
[15] validation_0-logloss:0.29127
```

```
[16] validation_0-logloss:0.28799
[17] validation_0-logloss:0.28873
[18] validation_0-logloss:0.28920
[19] validation_0-logloss:0.28884
[20] validation_0-logloss:0.28987
[21] validation_0-logloss:0.29086
[22] validation_0-logloss:0.29051
[23] validation_0-logloss:0.28783
[24] validation_0-logloss:0.28787
[25] validation_0-logloss:0.28888
[26] validation_0-logloss:0.28959
[27] validation_0-logloss:0.28966
[28] validation_0-logloss:0.29026
[29] validation_0-logloss:0.29129
[30] validation_0-logloss:0.29195
[31] validation_0-logloss:0.29206
[32] validation_0-logloss:0.29246
[33] validation_0-logloss:0.29040
[0] validation_0-logloss:0.52777
[1] validation_0-logloss:0.43750
[2] validation_0-logloss:0.38840
[3] validation_0-logloss:0.35246
[4] validation_0-logloss:0.32736
[5] validation_0-logloss:0.31123
[6] validation_0-logloss:0.30156
[7] validation_0-logloss:0.29543
[8] validation_0-logloss:0.29108
[9] validation_0-logloss:0.28816
[10] validation_0-logloss:0.29117
[11] validation_0-logloss:0.29301
[12] validation_0-logloss:0.29640
[13] validation_0-logloss:0.29625
[14] validation_0-logloss:0.29656
[15] validation_0-logloss:0.30112
[16] validation_0-logloss:0.30018
[17] validation_0-logloss:0.29921
[18] validation_0-logloss:0.30100
[0] validation_0-logloss:0.52703
[1] validation_0-logloss:0.43674
[2] validation_0-logloss:0.38237
[3] validation_0-logloss:0.34501
[4] validation_0-logloss:0.32304
[5] validation_0-logloss:0.30740
[6] validation_0-logloss:0.29668
[7] validation_0-logloss:0.29257
[8] validation_0-logloss:0.29151
[9] validation_0-logloss:0.28918
[10] validation_0-logloss:0.29037
[11] validation_0-logloss:0.29218
[12] validation_0-logloss:0.29063
[13] validation_0-logloss:0.29522
[14] validation_0-logloss:0.29485
[15] validation_0-logloss:0.29618
[16] validation_0-logloss:0.30019
[17] validation_0-logloss:0.30195
[18] validation_0-logloss:0.30447
[19] validation_0-logloss:0.30894
[0] validation_0-logloss:0.52703
[1] validation_0-logloss:0.43674
[2] validation_0-logloss:0.38237
[3] validation_0-logloss:0.34501
[4] validation_0-logloss:0.32304
[5] validation_0-logloss:0.30740
[6] validation_0-logloss:0.29668
```

```
[7] validation_0-logloss:0.29257
[8] validation_0-logloss:0.29168
[9] validation_0-logloss:0.29008
[10] validation_0-logloss:0.28863
[11] validation_0-logloss:0.29265
[12] validation_0-logloss:0.29404
[13] validation_0-logloss:0.29427
[14] validation_0-logloss:0.29844
[15] validation_0-logloss:0.29993
[16] validation_0-logloss:0.29865
[17] validation_0-logloss:0.30166
[18] validation_0-logloss:0.30280
[19] validation_0-logloss:0.30454
[0] validation_0-logloss:0.52703
[1] validation_0-logloss:0.43674
[2] validation_0-logloss:0.38237
[3] validation_0-logloss:0.34501
[4] validation_0-logloss:0.32304
[5] validation_0-logloss:0.30740
[6] validation_0-logloss:0.29668
[7] validation_0-logloss:0.29257
[8] validation_0-logloss:0.29168
[9] validation_0-logloss:0.29008
[10] validation_0-logloss:0.28863
[11] validation_0-logloss:0.29265
[12] validation_0-logloss:0.29404
[13] validation_0-logloss:0.29427
[14] validation_0-logloss:0.29844
[15] validation_0-logloss:0.29993
[16] validation_0-logloss:0.29865
[17] validation_0-logloss:0.30166
[18] validation_0-logloss:0.30280
[19] validation_0-logloss:0.30454
[20] validation_0-logloss:0.30455
[0] validation_0-logloss:0.52957
[1] validation_0-logloss:0.43944
[2] validation_0-logloss:0.38977
[3] validation_0-logloss:0.35359
[4] validation_0-logloss:0.32719
[5] validation_0-logloss:0.31539
[6] validation_0-logloss:0.30277
[7] validation_0-logloss:0.29750
[8] validation_0-logloss:0.29406
[9] validation_0-logloss:0.29429
[10] validation_0-logloss:0.29286
[11] validation_0-logloss:0.29258
[12] validation_0-logloss:0.29295
[13] validation_0-logloss:0.29271
[14] validation_0-logloss:0.29479
[15] validation_0-logloss:0.29615
[16] validation_0-logloss:0.29605
[17] validation_0-logloss:0.29656
[18] validation_0-logloss:0.29605
[19] validation_0-logloss:0.29830
[20] validation_0-logloss:0.30036
```



A learning rate of 0.1 and maximum depth of 2 will be chosen.

## Final XGBoost Model

```
In [48]: xgb_model = XGBClassifier(n_estimators = 500, learning_rate = 0.1, max_depth
xgb_model.fit(X_train, y_train_encoded, eval_set=[(X_test, y_test_encoded)])
```

```
[0]    validation_0-logloss:0.63768
[1]    validation_0-logloss:0.58984
[2]    validation_0-logloss:0.55032
[3]    validation_0-logloss:0.51793
[4]    validation_0-logloss:0.49032
[5]    validation_0-logloss:0.46623
[6]    validation_0-logloss:0.44569
[7]    validation_0-logloss:0.42838
[8]    validation_0-logloss:0.41388
[9]    validation_0-logloss:0.40094
[10]   validation_0-logloss:0.39106
[11]   validation_0-logloss:0.38106
[12]   validation_0-logloss:0.37201
[13]   validation_0-logloss:0.36577
[14]   validation_0-logloss:0.35889
[15]   validation_0-logloss:0.35374
[16]   validation_0-logloss:0.34795
[17]   validation_0-logloss:0.34181
[18]   validation_0-logloss:0.33680
[19]   validation_0-logloss:0.33331
[20]   validation_0-logloss:0.32946
[21]   validation_0-logloss:0.32474
[22]   validation_0-logloss:0.32158
[23]   validation_0-logloss:0.31762
[24]   validation_0-logloss:0.31474
```

```
[25] validation_0-logloss:0.31208
[26] validation_0-logloss:0.30892
[27] validation_0-logloss:0.30728
[28] validation_0-logloss:0.30577
[29] validation_0-logloss:0.30347
[30] validation_0-logloss:0.30225
[31] validation_0-logloss:0.30010
[32] validation_0-logloss:0.29872
[33] validation_0-logloss:0.29715
[34] validation_0-logloss:0.29587
[35] validation_0-logloss:0.29480
[36] validation_0-logloss:0.29413
[37] validation_0-logloss:0.29373
[38] validation_0-logloss:0.29266
[39] validation_0-logloss:0.29171
[40] validation_0-logloss:0.29106
[41] validation_0-logloss:0.29047
[42] validation_0-logloss:0.29003
[43] validation_0-logloss:0.28936
[44] validation_0-logloss:0.28903
[45] validation_0-logloss:0.28828
[46] validation_0-logloss:0.28838
[47] validation_0-logloss:0.28799
[48] validation_0-logloss:0.28719
[49] validation_0-logloss:0.28699
[50] validation_0-logloss:0.28654
[51] validation_0-logloss:0.28635
[52] validation_0-logloss:0.28575
[53] validation_0-logloss:0.28559
[54] validation_0-logloss:0.28547
[55] validation_0-logloss:0.28495
[56] validation_0-logloss:0.28515
[57] validation_0-logloss:0.28514
[58] validation_0-logloss:0.28502
[59] validation_0-logloss:0.28507
[60] validation_0-logloss:0.28473
[61] validation_0-logloss:0.28470
[62] validation_0-logloss:0.28547
[63] validation_0-logloss:0.28556
[64] validation_0-logloss:0.28537
[65] validation_0-logloss:0.28519
[66] validation_0-logloss:0.28461
[67] validation_0-logloss:0.28474
[68] validation_0-logloss:0.28461
[69] validation_0-logloss:0.28509
[70] validation_0-logloss:0.28488
[71] validation_0-logloss:0.28504
[72] validation_0-logloss:0.28507
[73] validation_0-logloss:0.28469
[74] validation_0-logloss:0.28458
[75] validation_0-logloss:0.28510
[76] validation_0-logloss:0.28464
[77] validation_0-logloss:0.28419
[78] validation_0-logloss:0.28462
[79] validation_0-logloss:0.28435
[80] validation_0-logloss:0.28529
[81] validation_0-logloss:0.28499
[82] validation_0-logloss:0.28516
[83] validation_0-logloss:0.28499
[84] validation_0-logloss:0.28521
[85] validation_0-logloss:0.28496
[86] validation_0-logloss:0.28458
[87] validation_0-logloss:0.28536
```

Out[48]:

```
▼ XGBClassifier
  XGBClassifier(base_score=None, booster=None, callbacks=None,
                 colsample_bylevel=None, colsample_bynode=None,
                 colsample_bytree=None, device=None, early_stopping_roun
ds=10,
                 enable_categorical=False, eval_metric=None, feature_typ
es=None,
                 gamma=None, grow_policy=None, importance_type=None,
                 interaction_constraints=None, learning_rate=0.1, max_bi
n=None,
                 max_cat_threshold=None, max_cat_to_onehot=None,
                 max_delta_step=None, max_depth=2, max_leaves=None,
```

## Best Iteration

In [49]:

```
xgb_model.best_iteration
```

Out[49]:

Although set the number of trees to be 500, but actually the best number is 77.

## Evaluation

In [50]:

```
y_pred = xgb_model.predict(X_test)
```

## Classification Report

In [51]:


```
print(classification_report(y_test_encoded, y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.91	0.89	251
1	0.91	0.86	0.88	249
accuracy			0.89	500
macro avg	0.89	0.89	0.89	500
weighted avg	0.89	0.89	0.89	500

## Confusion Matrix

In [52]:


```
cm = confusion_matrix(y_test_encoded, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix = cm, display_labels = class_l
disp.plot(cmap='Blues')
plt.title("XGBoost Confusion Matrix")
plt.show()
```



## Importance

In [53]:

```
importances = xgb_model.feature_importances_
feat_importances = pd.Series(importances, index = features)
feat_importances.sort_values(ascending = False).plot(kind = 'bar')
plt.title("Feature Importance")
plt.show()
```



## 5. Conclusion

- Random Forest performs slightly better on this dataset in terms of overall accuracy (0.90) and F1-score (Çerçevelek: 0.90, Ürgüp Sivrisi: 0.89).
- XGBoost is also showing good F1-score (Çerçevelek: 0.89, Ürgüp Sivrisi: 0.88) and accuracy (0.89).
- Both models are strong classifiers.
- Top three important features for Random Forest are Eccentricity, Compactness, and Aspect Ration.
- Top three important features for XGBoost are Aspect Ration, Compactness, and Eccentricity.
- Top features are the same except their ranking.